

CONTACT

paddy.iyer@gmail.com
[linkedin.com/in/paddyiyer](https://www.linkedin.com/in/paddyiyer)
paddyspeaks.com
github.com/paddyriyer
San Francisco Bay Area

TECHNICAL SKILLS

Spark Hive Hadoop Databricks Delta
Lake Snowflake Redshift dbt
Airflow Dataswarm AWS Azure
GCP Terraform Kubernetes Python
SQL Scala LLM/RAG Pipelines Claude
Code AI Agents Prompt Eng.
LangChain GitHub Copilot MLflow
Diff. Privacy Unity Catalog Data Mesh

EDUCATION

Data Engineering Cert.
UC Santa Cruz, 2013–14
B.S. Electrical Engineering
Sardar Patel College of Eng.
Diploma, Electronics & Comms
Technical Board, Tamil Nadu

PUBLICATIONS

Cloud Computing
All About Big Data
Future Trends in BI

AI & THOUGHT LEADERSHIP

160+ articles on paddyspeaks.com
Claude Code skill development
Vibe coding methodology
AI agent architecture writing

OPEN-SOURCE & COMMUNITY

Interview Studio — free Data & AI
Engineering prep
22 data-model deep-dives +
1,500+ questions
In-browser SQL / Python (WASM),
no backend

EARLIER CAREER

Data Engineer Consultant

LinkedIn
Dec 2017 – Apr 2018
GDPR compliance — PII encryption into
Dali storage (Hive, Python, Pig)

Lead Data Engineer

Meta
Sep 2015 – Sep 2017
80+ Dataswarm pipelines for petabyte-
scale ads; Cross Device Insights, Outcomes
Datamart + Norms DB

Data Engineer

GREE International
2014 – 2015
PII masking, Vertica→Redshift, 1,000+
table optimization

Data Architect

Chegg Inc.
Oct 2013 – Apr 2014

Technical Director

Model N
2011 – 2013
Life Sciences BI, cloud migration, ETL

Paddy Iyer

DATA ENGINEERING LEADER • STRATEGIC DATA ARCHITECT • AI & PRIVACY ENGINEERING • AI INNOVATOR

15+ years driving enterprise data transformation for global tech leaders across cloud/SaaS, fintech, gaming, and consumer platforms. Expert in cloud-native architectures, real-time analytics, privacy-first pipelines, and AI-augmented data operations. Proven builder of high-trust, high-performance data teams at petabyte scale. Published writer and thought leader on data architecture, AI agents, and the intersection of ancient wisdom with modern technology.

Modern Data Architecture Privacy-First Design Consent Architecture AI Agents &
LLM Pipelines Prompt Engineering Data Governance (C1–C5) Petabyte-Scale
Systems Global Team Leadership Data Storytelling

FEATURED EXPERIENCE

Data Engineer — Privacy & Consent InfrastructureApr 2024 –

Present

Meta

- Landed 1,000+ production changes across a multi-year privacy migration — sustained solo throughput normally requiring a team of 3–4 engineers
- Cleared the final blocker to a company-wide data warehouse consolidation program. My workstream covered the last ~20% of warehouse data — the portion nobody else could touch because it sat directly in the advertising revenue path. Completing it took the program from 80% to 100% coverage and enabled its formal close-out
- Executed that migration with zero revenue impact on the highest-sensitivity assets in the warehouse, under a constraint that any error would surface as ads revenue loss rather than a test failure
- Proved the pre-deployment validation stack was structurally blind — permission checks returned PASS for tables that didn't exist and never modeled runtime enforcement. Built the hard-gate guard system that closed the gap; every gate traces to a specific postmortem
- Reduced an entire failure class to an offline-checkable invariant. Verified against the full asset inventory: 81% of migrations required a namespace change, not the name change the tooling assumed
- Owned response for 5 production incidents including 2 SEV3s (one at \$5.27M/day) — fix-forward, zero reverts, permission grants coordinated across 6+ owning teams
- Co-built an AI coding skill encoding the full remediation methodology: engineer ramp-up 2 weeks → 2 days, per-asset cycle time down 3–5×
- Work recognized in a company-wide announcement to the Data organization, with the difficulty of this workstream called out explicitly by senior engineering leadership

EARLIER CAREER

Lead Data Architect — Partner Data EngineeringJan 2022 – Jan 2024

VMware

Strategic architect for VMware's Partner Data Platform — unifying fragmented partner data into a single governed analytics layer

- Designed and built the Partner Data Platform from the ground up, unifying partner data from 10+ disparate sources into a single governed analytics layer — eliminating silos that had blocked cross-functional decision-making for years
- Implemented VMware's first structured data governance for partner data (classification, lineage, policy enforcement) plus a self-service data catalog — cutting tribal-knowledge reliance by 60% and ad-hoc request volume by ~40%, and achieving audit-ready status
- Pioneered an AI copilot for retrieval optimization, coding, and governance — one of VMware's earliest LLM-assisted data engineering deployments, reducing routine query development time by ~30%
- Achieved 40%+ Spark processing improvements via broadcast joins, caching, and skew management — bringing SLA-critical reports inside their delivery windows for the first time

modularization

Architect

CallidusCloud

2005 - 2011

Incentive comp analytics, BusinessObjects XI

DW Architect

Hewlett Packard

2001 - 2005

Enterprise DW + 8 datamarts; ETL 18→3 hrs; 40-50% sales lift via clickstream

- Engaged director/VP-level stakeholders to define a unified data vision, securing budget and headcount that grew the team from 3 to 8 engineers during the engagement

Senior Consultant — Ads, Commerce & Privacy *Jul 2019 - Dec 2021* Meta

Led high-impact projects across Meta's Ads, Commerce, and Privacy teams, architecting secure, scalable, and compliant data platforms

- Led privacy remediation across 1,200+ SMB 2.0, Customer Journey, and BPO tables — resolving critical compliance-risk design flaws and establishing remediation patterns later adopted at scale in the Safe Ads program
- Spearheaded migration of 1,000+ Hive pipelines to Spark, automating ~70% of conversion tasks — compressing a projected 12-month migration to under 6 months
- Owned the SMB funnel redesign, improving query performance 3x and scaling the pipeline from thousands to tens of thousands of small-business advertisers without infrastructure changes
- Enhanced Hive anonymization processes with privacy engineers, enabling ads and commerce data products to pass privacy audits without manual intervention — cutting audit preparation effort by ~60%