

CONTACT

paddy.iyer@gmail.com
linkedin.com/in/paddyiyer
paddyspeaks.com
github.com/paddyriyer
San Francisco Bay Area

TECHNICAL SKILLS

Spark Hive Hadoop Databricks Delta
Lake Snowflake Redshift dbt Airflow
Dataswarm AWS Azure GCP
Terraform Kubernetes Python SQL
Scala LLM/RAG Pipelines Claude Code
AI Agents Prompt Eng. LangChain
GitHub Copilot MLflow Diff. Privacy
Unity Catalog Data Mesh

EDUCATION

Data Engineering Cert.

UC Santa Cruz, 2013–14

B.S. Electrical Engineering

Sardar Patel College of Eng.

Diploma, Electronics & Comms

Technical Board, Tamil Nadu

PUBLICATIONS

- Cloud Computing
- All About Big Data
- Future Trends in BI

AI & THOUGHT LEADERSHIP

- 160+ articles on paddyspeaks.com
- Claude Code skill development
- Vibe coding methodology
- AI agent architecture writing

OPEN-SOURCE & COMMUNITY

- Interview Studio — free Data & AI Engineering prep
- 22 data-model deep-dives + 1,500+ questions
- In-browser SQL / Python (WASM), no backend

EARLIER CAREER

Data Engineer Consultant

[LinkedIn](#)

Dec 2017 – Apr 2018

GDPR compliance — PII encryption into Dali storage (Hive, Python, Pig)

Lead Data Engineer

[Meta](#)

Sep 2015 – Sep 2017

80+ Dataswarm pipelines for petabyte-scale ads; Cross Device Insights, Outcomes Datamart + Norms DB

Data Engineer

[GREE International](#)

2014 – 2015

PII masking, Vertica→Redshift, 1,000+ table optimization

Data Architect

[Chegg Inc.](#)

Oct 2013 – Apr 2014

Technical Director

[Model N](#)

2011 – 2013

Life Sciences BI, cloud migration, ETL modularization

Architect

Paddy Iyer

DATA ENGINEERING LEADER • STRATEGIC DATA ARCHITECT • AI & PRIVACY ENGINEERING • AI INNOVATOR

15+ years driving enterprise data transformation for global tech leaders across cloud/SaaS, fintech, gaming, and consumer platforms. Expert in cloud-native architectures, real-time analytics, privacy-first pipelines, and AI-augmented data operations. Proven builder of high-trust, high-performance data teams at petabyte scale. Published writer and thought leader on data architecture, AI agents, and the intersection of ancient wisdom with modern technology.

Modern Data Architecture Privacy-First Design Consent Architecture AI Agents & LLM Pipelines Prompt Engineering Data Governance (C1–C5) Petabyte-Scale Systems Global Team Leadership Data Storytelling

FEATURED EXPERIENCE

Data Engineer — Privacy & Consent Infrastructure Apr 2024 – Present [Meta](#)

Drove the data engineering execution of a company-critical migration moving 1,100+ advertising data assets from broad-access to consent-first architecture. Landed 800+ changes across four privacy domains. What I actually got paid for: finding the failures nobody's tooling could see

Finding the Failures the Tooling Couldn't See

- Proved the pre-deployment validation stack was structurally blind — per-table permission checks returned PASS for tables that didn't exist, and never modeled the query-level enforcement that actually runs in production. Diffs cleared every gate, then broke
- Reduced a class of migration failures to a mechanical, offline-checkable invariant: the source namespace was encoded in the old table name all along. Verified against the full asset inventory — 81% of migrations required a namespace change, not a name change. Made the correct transform regex-provable instead of lookup-dependent
- Built the hard-gate guard system that turned each production incident into a pre-land block. Every gate in the playbook traces to a specific postmortem
- Separated silent failure from loud failure — showed that the “obvious fix” for a resolution error swapped a paging error for a silently skipped subquery. Blocked it as strictly worse, and documented why
- Enforced scope integrity across a large migration batch: privacy annotations treated as invariants, not redundant syntax. Blocked collateral edits that removed safety controls to satisfy a linter

Force Multiplication

- Co-built an AI coding skill encoding the full remediation methodology — new engineers productive in 2 days instead of 2 weeks; per-asset cycle time down 3–5×
- Sustained 800+ landed changes solo — throughput normally requiring 3–4 engineers

Incident Leadership

- Owned response for 5 production failures including 2 SEV3s (one at \$5.27M/day). Fix-forward, zero reverts, grants coordinated across 6+ owning teams. Director-level escalations closed with written RCA
- Triage rule I enforced: reason from the runtime error, never from a plausible mechanism. When your analysis contradicts the pipeline owner's, you're missing a layer

EARLIER CAREER

Lead Data Architect — Partner Data Engineering Jan 2022 – Jan 2024 [VMware](#)

Strategic architect for VMware's Partner Data Platform — unifying fragmented partner data into a single governed analytics layer

- Designed and built the Partner Data Platform from the ground up, unifying partner data from 10+ disparate sources into a single governed analytics layer — eliminating silos that had blocked cross-functional decision-making for years
- Implemented VMware's first structured data governance for partner data (classification, lineage, policy enforcement) plus a self-service data catalog — cutting tribal-knowledge reliance by 60% and ad-hoc request volume by ~40%, and achieving audit-ready status

CallidusCloud

2005 – 2011

Incentive comp analytics, BusinessObjects XI

DW Architect

Hewlett Packard

2001 – 2005

Enterprise DW + 8 datamarts; ETL 18→3 hrs;
40–50% sales lift via clickstream

- Pioneered an AI copilot for retrieval optimization, coding, and governance — one of VMware's earliest LLM-assisted data engineering deployments, reducing routine query development time by ~30%
- Achieved 40%+ Spark processing improvements via broadcast joins, caching, and skew management — bringing SLA-critical reports inside their delivery windows for the first time
- Engaged director/VP-level stakeholders to define a unified data vision, securing budget and headcount that grew the team from 3 to 8 engineers during the engagement

Senior Consultant — Ads, Commerce & Privacy *Jul 2019 – Dec 2021*

Meta

Led high-impact projects across Meta's Ads, Commerce, and Privacy teams, architecting secure, scalable, and compliant data platforms

- Led privacy remediation across 1,200+ SMB 2.0, Customer Journey, and BPO tables — resolving critical compliance-risk design flaws and establishing remediation patterns later adopted at scale in the Safe Ads program
- Spearheaded migration of 1,000+ Hive pipelines to Spark, automating ~70% of conversion tasks — compressing a projected 12-month migration to under 6 months
- Owned the SMB funnel redesign, improving query performance 3x and scaling the pipeline from thousands to tens of thousands of small-business advertisers without infrastructure changes
- Enhanced Hive anonymization processes with privacy engineers, enabling ads and commerce data products to pass privacy audits without manual intervention — cutting audit preparation effort by ~60%